

The Halting Problem

Reading

Section 4.2

The Universal Turing Machine

The work we did in the previous section naturally extends to Turing Machines. In fact we can define:

$$A_{\text{TM}} = \{\langle M, w \rangle \mid M \text{ is a TM that accepts } w\}$$

So this language captures all pairs of a Turing Machine and a string, such that the Turing Machine accepts the string.

A_{TM} is Turing-recognizable.

The machine that recognizes this language is called a *Universal Turing Machine*: It is effectively a Turing Machine that simulates the computation of all Turing Machines. It goes as follows:

1. On input a TM M and a string w .
2. It simulates M on input that string w .
3. If M enters its accept or reject states, then our universal Turing machine also enters its accept (resp. reject) state.
4. If M loops (runs forever), so does our universal Turing machine.

Notice that this is most definitely not a decider: It has no way to determine whether M will loop or not, other than actually simulating it and running forever.

The Halting Problem

The Halting Problem effectively boils down to the following:

A_{TM} is not decidable.

In other words, there can never be a Turing Machine/program that is guaranteed to terminate and that can decide, given an arbitrary Turing Machine/program and its input, whether that program would terminate. In terms of programming languages, this means that we cannot write a program that can tell us if programs in our language

would terminate. This has implications in terms of what a type system can do, which go a bit beyond the scope of this course.

This is also our first example of a specific language that is Turing-recognizable but not decidable.

Now let us discuss the proof of this extremely important result. It is based on another assertion:

Consider the language:

$$\text{NONSELFACC} = \{\langle M \rangle \mid M \text{ does not accept on input } \langle M \rangle\}$$

of all Turing Machines that do not accept when given themselves as input.

Then NONSELFACC is not decidable.

Before we prove this result, let us see why this would imply that A_{TM} is also not decidable.

If A_{TM} were decidable, denote by H a decider. We then construct a decider for NONSELFACC as follows:

- Given a Turing machine M , use H to run M on input its string representation $\langle M \rangle$.
- If H accepts, then we reject.
- If H rejects, then we accept.

This describes a decider for NONSELFACC. But since that was not a decidable language, we have a contradiction.

Now let us discuss why NONSELFACC is not decidable. Suppose instead that it was decidable, and that D was its decider. Then we have a problem:

- If $D \in \text{NONSELFACC}$, then it means that NONSELFACC's decider must have accepted D , and so running D on input $\langle D \rangle$ must have accepted. But since D was a decider for NONSELFACC, this means that $D \notin \text{NONSELFACC}$.
- Conversely, if $D \notin \text{NONSELFACC}$, then this means that running D with input itself, $\langle D \rangle$, would accept. But since D is the decider of NONSELFACC and it just accepted input $\langle D \rangle$, this must mean that $D \in \text{NONSELFACC}$.

So we have a contradiction, as D cannot both be in NONSELFACC and not be in it. So D cannot exist in the first place.

A non-recognizable language

We end this chapter with an example of a language that's not even Turing-recognizable.

The language that is the complement of A_{TM} is not Turing-recognizable.

In general, the complement of any language that is Turing-recognizable but not decidable, is itself not Turing-recognizable.

This follows simply from the observation that if both a language and its complement were Turing-recognizable, then we could build a decider for the language:

1. On input w :
2. Simultaneously simulate the recognizer Turing Machines for the language and its complement, both on input w .
3. Since w is either in the language or its complement, one of those recognizers will have to accept w . We can then terminate with either accept or reject, depending on which recognizer accepted.

Exercise: What about the following slightly different language:

Show that the language

$$L = \{\langle M, w \rangle \mid M \text{ is a TM that does not accept on input } w\}$$

is not Turing-recognizable.