

Turing Machines

Reading

Section 3.1, 3.2

Practice problems (page 159): 3.5, 3.8, 3.15, 3.22

Challenge: 3.19, 3.20

Turing Machines

Turing Machines (TM) are a computational structure that is quite a step up from pushdown automata in terms of capabilities. In fact in a certain sense they are strong enough to capture the idea of computation in general.

Here are some comparisons between Turing Machines and PDAs/DFAs:

- There is no separate input and stack space. There is a single infinitely long “tape”, which is initially populated by the input string.
- The TM can freely move along the tape in either direction. It is not restricted by input order and stack discipline.
- Like DFAs and PDAs, the TM has a finite set of states that dictate its options.
- The TM’s transition function is deterministic. We will in fact prove that adding nondeterminism does not increase the TM’s computational power.
- The TM has special “accept” and “reject” states, and the computation ends the moment those are entered.
- The computation can go on forever. For CFGs in Chomsky Normal form there are bounds to the number of steps needed to reach a string of a given length. Nothing equivalent holds for TMs.

Turing Machines operate on an infinitely long tape, and at any given time the machine is pointing at some location in the tape, and is at some state. To transition, it can make a move left or right.

Here is the formal definition:

Turing Machines

A **Turing Machine** is a 7-tuple $(Q, \Sigma, \Gamma, \delta, q_0, q_{\text{accept}}, q_{\text{reject}})$ where:

1. Q is a finite nonempty set of states.
2. Σ is a finite **input alphabet**, not containing a special “blank symbol”.
3. Γ is a finite **tape alphabet**, containing $\Sigma \subset \Gamma$ and a special **blank symbol** \sqcup representing the lack of a value at that location in the “tape”.

4. $\delta: Q \times \Gamma \rightarrow Q \times \Gamma \times \{L, R\}$ is a transition function. Based on the value at the current location in the tape, as well as the current state, the automaton can change the value at the location, switch to a new state, and also move the pointer/head to the left or to the right (L/R).
5. $q_0 \in Q$ is a start state.
6. $q_{\text{accept}} \in Q$ is the accept state
7. $q_{\text{reject}} \in Q$ is the reject state, *distinct* from q_{accept} .

Computation in a Turing Machine

- The initial input to the machine is placed at the leftmost n squares of the tape. The rest of the tape is filled with the blank symbol.
- The TM starts at the state q_0 , and with the marker pointing at the leftmost location.
- The TM follows the rules dictated by the transition function δ , until it reaches the accept or reject states.
- If the TM ever tries to move to the left of the leftmost square, it just stays in that square.

We often describe Turing machines more informally than listing states and transitions, by describing the steps to be performed, with an understanding that if needed we could write down states and transitions to carry out those steps. Let us look at an example of such an informal description:

Design a TM that recognizes the language $L = \{w\#w \mid w \in \{0,1\}^*\}$.

You will recall that this was one of the languages that pushdown automata could not handle.

Here's how this TM could go:

1. We assume that our tape alphabet Γ contains “dotted” versions of our symbols, $\dot{0}, \dot{1}$. We can use those to mark locations.
2. At the beginning we are the leftmost element. We replace it with a dotted version of it, and move to a state that remembers what that value was and searches for the $\#$ marker.
3. After we reach the marker, we look at the location after it and whether it matched the value from the other side, that we were keeping track of. If it does not, then we move to reject state. If it does, we mark the location, either via a dotted version of the value or via an x.
4. Now we move left till we find the $\#$ marker, then keep moving till we find the dotted value. We move to the value after it, (i.e. the first non-dotted value), dot that and remember it, then move to find its partner on the right side.
5. We continue this back and forth until the point that we have looked at all the elements before the $\#$ marker. This means we have matched everything to the left of the $\#$ marker with the corresponding spots on the right.

6. All that is left is to travel to the right past all the x-ed out values, and see if the next thing is the blank, or if there are still more values out there. If it is blank, it means we have matched the string before # with the string after it, and so we move to accept. If not, we move to reject.

We could imagine using a dozen or so different states to carry out the above steps.

Let us try a variation:

Design a TM that recognizes the language $L = \{ww \mid w \in \{0,1\}^*\}$.

In this instance we don't have the marker to separate the two words. We will need to find the middle of the string ourselves.

1. Start at the leftmost element, dot it, and start moving right till you find the other end.
2. Dot that element, and move left till you find that previously dotted element. dot the one after it.
3. Keep going back and forth marking the elements you have seen, one on each end.
4. When we reach the middle area, we can detect if the string had odd length: We just dotted the middle element as belonging to the left half of the string, and when we go right to find an undotted element there isn't one remaining. We can transition to a reject state right away in this case.
5. Otherwise we have determined the string has even length, and we are standing at the beginning of the right half of it. We can now proceed in two ways:
6. One way is to undot that element and remember what it is, then move all the way at the beginning of the string (should have dotted that first element in a more special way) and see if it matches, and if it does cross it out. Then take the next element, cross it out and remember it and look for it on the second half of the string, and so on.
7. Another alternative is to insert a new marker at that middle position. We can then "shift" the second half of the string one square to the right, as follows:
 - Place the marker, and remember what was there before, move to the right.
 - Place in that square the element you are remembering, and remember the new element, move to the right.
 - Keep doing that until you have arrived at the blank symbol, marking the end of the string. You then insert that last element, and transition to whatever you need to do next.
 - Now our string has obtained the form described in the previous problem, so we can use the same method we used for that one.

This idea of reusing a previously constructed TM as a helper is a key tool in the methodology of Turing Machines.

Before we move on to more examples, there is one important distinction we need to make.

A TM on a given input can result in one of three outcomes: *accept*, *reject* or *loop*. A **loop** means that the machine goes on forever, and never reaches an accept or reject state. Unlike with DFAs and PDAs, it is not at all easy to detect the cases where the TM will loop.

We say that a TM **recognizes** a language L , if the TM results in *accept* for all strings in the language, and not in accept for all strings not in the language. So it may be the case that the TM results in a loop for some strings not in the language, and that is OK. We say that the language L is **Turing-recognizable**, if there is a TM recognizing it.

We say that a TM **decides** a language L , if the TM results in *accept* for all strings in the language and in *reject* for all strings not in the language. In other words, the TM never loops. Such a TM is called a **decider**. We call a language L **Turing-decidable**, or just **decidable**, if some Turing machine decides it.

It should be clear that every decidable language is recognizable. The interesting fact is that the converse is not true: There are languages that are recognizable, but not decidable. The examples we have seen above are examples of decidable languages.

Practice: Describe TMs that decide the following languages:

1. $A = \{a^n b^n c^n \mid n \geq 0\}$
2. $B = \{a^i b^j c^k \mid 0 \leq i \leq j \leq k\}$
3. $C = \{a^i b^j \mid i \text{ is a positive multiple of } j\}$
4. $D = \{w\#t \mid w, t \in \{a, b\}^*, w \text{ is a substring of } t\}$
5. $E = \{w \in \{0, 1\}^* \mid \text{equal number of 0s, 1s}\}$

Turing Machine Variants

The remarkable fact about Turing Machines is that they are remarkably robust: Any variant we try turns out to have the same computational power as Turing Machines.

Stay option

A first simple variant is to allow three possible “moves”, instead of just left and right. There is also a S move, for staying at the same location.

This of course isn’t a considerable change: We can replace any “stay” move with a combination of a right move and a left move.

The essence of showing a variant is equivalent with the basic definition is to simulate one with the other.

Multitape machines

A more adventurous variant is one where we use multiple tapes. Each tape has its own head for reading and writing. Initially the input is all in the first tape, and the remaining tapes are blank. The transition function would depend on the values at all heads (at all the tapes), and will update all tapes at once:

$$\delta: Q \times \Gamma^k \rightarrow Q \times \Gamma^k \times \{L, R, S\}^k$$

It is clear that this notion is at least as strong as that of normal TMs. We will in fact show that multitape TMs are not stronger than normal TMs.

Every multitape Turing Machine has an equivalent single-tape Turing Machine.

The construction is instructive. The main idea is for our single-tape TM to emulate the operations of the multi-tape TM:

- We will store the contents of all the tapes in the multi-tape TM into a single tape.
 - Since at any given time only finitely many entries in the tape can be non-blank, this is possible.
 - Need a new symbol (say #) to separate the tapes.
 - Initial contents would look like #100110101#□#□#□#□#□#
- We use specially dotted versions of symbols to keep track of the heads in each tape of the multi-tape machine. So initially the tape starts with #100110101#□#□#□#□#□#
- In order to take a “step”:
 - Go through and read the dotted items. Finitely many possibilities, so we can keep track of all the combinations via different states.
 - Once all dotted items are read, use the multitape transition function to determine what to do next.
 - Based on that multitape transition function, go through our single tape and update the values on the dotted items, and also move the dots as dictated by the L/R values from the transition.
- If the multi-tape TM enters an accept/reject state, then so does the single-tape TM.
- Otherwise it continues to emulate the multi-tape TM.

Non-deterministic Turing Machines

A non-deterministic Turing Machine is, as with previous notions of non-determinism, one whose transitions contain multiple alternatives. So the transition function looks like:

$$\delta: Q \times \Gamma \rightarrow \mathcal{P}(Q \times \Gamma \times \{L, R\})$$

Such a Turing Machine accepts if any computation branch results in an accept state. It does not reject when it meets a reject state. It will only reject if it has exhausted all computation paths without reaching an accept state.

It may appear as if this variant will give us more power, but that turns out not to be the case:

Every non-deterministic Turing Machine has an equivalent deterministic Turing Machine

We will instead show that a non-deterministic Turing Machine is equivalent to a multi-tape Turing Machine with 3 tapes.

First some preparations. We can make the non-determinism a bit more concrete. Note that the set $Q \times \Gamma$ is finite in size, and for each of those “states of the TM” the corresponding options provided by the transition function are finitely many. We fix an order for these options, and number them. This way we can imagine a sequence of numbers like 13211 to suggest a specific path of computation: “From our start configuration we chose the first of the available options, then from the options we had at that state we chose the third one, and from the options we had at that location we chose the second one, and so on”. We can imagine all these options represented by a tree, whose nodes represent steps in the computation, and where each node has children corresponding to the different options provided by the transition function δ for the configuration described by that node.

Effectively what we want to do is explore all paths of this “computation tree”, in search of an accept or reject result. We have to be careful however: A depth-first traversal of the tree might get us down an infinite path from which we might never escape, even though an accept path might have existed down another direction in the tree. For that reason we take a breadth first approach: We take only one step at each of the possible states we might be in, resulting in a new set of possible states, before moving deeper in. Essentially we try all ways of say doing 3 steps from the start state, before we try any of the ways of doing 4 steps.

Now we are ready to describe the construction in more detail. The construction involves 3 tapes:

- The first tape contains the unaltered input to the TM. It uses the alphabet Σ of the language we are working with.
- The second tape is used to carry out a “emulation” of a particular path down the non-deterministic TM’s computation tree. It uses the alphabet Γ of the non-deterministic machine we are trying to emulate.
- The third tape is used to keep track of the path in the tree that we are currently exploring.

The computation goes as follows:

1. “Advance” the path in the third tape.

- Increase the last number in the tape, representing the last leg of the previous computation. Since we just finished exploring that “child”, we want to now move to its sibling.
- If that number has exceeded the number of children at that point, reset it to 1 then move to the previous entry in the tape (its parent) and advance that, looking at *its* sibling. Keep moving up the tree (earlier in the tape) as needed.
- If we end up increasing the first number in the tape past the number of children at that level, this means that we have exhausted all trees of the given depth, and we need to increase the depth: Move to the first non-blank entry in the tape and put 1 in.

2. “Emulate” the path currently indicated by the third tape.

- Move the head of the third tape at the beginning
- Clear the second tape, and write the input from the first tape in it. Place the head of the second tape at the beginning
- Following the (finite) path described by the third tape, and the instructions at the corresponding configurations of the non-deterministic machine, emulate on the second tape the work that the non-deterministic machine would have done.
- If you run into the accept state of the non-deterministic machine, then report accepting. If we run into a reject state, treat it as a loop and move on to the next path.

3. Go back to 1 and repeat

So essentially every time the deterministic TM wants to take a step, it has to reset to the nondeterministic TM to the “beginning”, and repeat all the steps that had been taken up to that point, then follow them with one more step. Then resets back to the beginning to explore another option, and so on and so forth. It is a much slower process, but we are not currently concerned with speed, only with feasibility.

So this way, if the non-deterministic TM would have reached an accept state, it would have done so in a finite number of steps, say n steps, so it would have corresponded to some path of depth n in the computation tree. This is a path that the deterministic TM will explore in one of its first 2^n loops through the steps 1-3 above.

This essentially completes the proof.

From this we have some consequences:

A language is Turing-recognizable if and only if there is a non-deterministic Turing Machine that recognizes it.

Decidability is a bit trickier:

A non-deterministic TM is called a decider if all branches of the computation tree halt on all inputs. It is then said to “accept” an input string if there is at least one branch that ends in the accept state.

Consider how you would modify the above construction in order to prove the following:

A non-deterministic TM that is a decider has an equivalent deterministic TM that is a decider.

A language is Turing-decidable if and only if there is a non-deterministic Turing Machine that decides it.